

## ✓ Handwritten digits recognition (using Convolutional Neural Network)

```
# Selecting Tensorflow version v2 (the command is relevant for Colab only).
%tensorflow_version 2.x
```

↗ Colab only includes TensorFlow 2.x; %tensorflow\_version has no effect.

```
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sn
import numpy as np
import pandas as pd
import math
import datetime
import platform

print('Python version:', platform.python_version())
print('Tensorflow version:', tf.__version__)
print('Keras version:', tf.keras.__version__)
```

↗ Python version: 3.11.12  
Tensorflow version: 2.18.0  
Keras version: 3.8.0

## ✓ Configuring Tensorboard

We will use [Tensorboard](#) to debug the model later.

```
# Load the TensorBoard notebook extension.
# %reload_ext tensorboard
%load_ext tensorboard
```

```
# Clear any logs from previous runs.
!rm -rf ./logs/
```

## ✓ Load the data

The **training** dataset consists of 60000 28x28px images of hand-written digits from 0 to 9.

The **test** dataset consists of 10000 28x28px images.

```
mnist_dataset = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist_dataset.load_data()
```

↗ Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 ————— 0s 0us/step

```
print('x_train:', x_train.shape)
print('y_train:', y_train.shape)
print('x_test:', x_test.shape)
print('y_test:', y_test.shape)
```

↗ x\_train: (60000, 28, 28)  
y\_train: (60000,)  
x\_test: (10000, 28, 28)  
y\_test: (10000,)

```
# Save image parameters to the constants that we will use later for data re-shaping and for model training.
(_, IMAGE_WIDTH, IMAGE_HEIGHT) = x_train.shape
IMAGE_CHANNELS = 1
```

```
print('IMAGE_WIDTH:', IMAGE_WIDTH);
print('IMAGE_HEIGHT:', IMAGE_HEIGHT);
print('IMAGE_CHANNELS:', IMAGE_CHANNELS);
```




↗ IMAGE\_WIDTH: 28  
IMAGE\_HEIGHT: 28

IMAGE\_CHANNELS: 1

Explore the data

Here is how each image in the dataset looks like. It is a 28x28 matrix of integers (from 0 to 255 ). Each integer represents a color of a pixel.

```
pd.DataFrame(x_train[0])
```




	0	1	2	3	4	5	6	7	8	9	...	18	19	20	21	22	23	24	25	26	27
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0	0	0	0	...	175	26	166	255	247	127	0	0	0	0
6	0	0	0	0	0	0	0	0	30	36	...	225	172	253	242	195	64	0	0	0	0
7	0	0	0	0	0	0	0	49	238	253	...	93	82	82	56	39	0	0	0	0	0
8	0	0	0	0	0	0	0	18	219	253	...	0	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	80	156	...	0	0	0	0	0	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0	14	...	0	0	0	0	0	0	0	0	0	0
11	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
12	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
13	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
14	0	0	0	0	0	0	0	0	0	0	...	25	0	0	0	0	0	0	0	0	0
15	0	0	0	0	0	0	0	0	0	0	...	150	27	0	0	0	0	0	0	0	0
16	0	0	0	0	0	0	0	0	0	0	...	253	187	0	0	0	0	0	0	0	0
17	0	0	0	0	0	0	0	0	0	0	...	253	249	64	0	0	0	0	0	0	0
18	0	0	0	0	0	0	0	0	0	0	...	253	207	2	0	0	0	0	0	0	0
19	0	0	0	0	0	0	0	0	0	0	...	250	182	0	0	0	0	0	0	0	0
20	0	0	0	0	0	0	0	0	0	0	...	78	0	0	0	0	0	0	0	0	0
21	0	0	0	0	0	0	0	0	23	66	...	0	0	0	0	0	0	0	0	0	0
22	0	0	0	0	0	0	18	171	219	253	...	0	0	0	0	0	0	0	0	0	0
23	0	0	0	0	55	172	226	253	253	253	...	0	0	0	0	0	0	0	0	0	0
24	0	0	0	0	136	253	253	253	212	135	...	0	0	0	0	0	0	0	0	0	0
25	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
26	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
27	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

28 rows × 28 columns


This matrix of numbers may be drawn as follows:

```
plt.imshow(x_train[0], cmap=plt.cm.binary)
plt.show()
```



Let's print some more training examples to get the feeling of how the digits were written.

```
numbers_to_display = 25
num_cells = math.ceil(math.sqrt(numbers_to_display))
plt.figure(figsize=(10,10))
for i in range(numbers_to_display):
    plt.subplot(num_cells, num_cells, i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i], cmap=plt.cm.binary)
    plt.xlabel(y_train[i])
plt.show()
```



## ✓ Reshaping the data

In order to use convolution layers we need to reshape our data and add a color channel to it. As you've noticed currently every digit has a shape of (28, 28) which means that it is a 28x28 matrix of color values from 0 to 255. We need to reshape it to (28, 28, 1) shape so that each pixel potentially may have multiple channels (like Red, Green and Blue).

```
x_train_with_channels = x_train.reshape(
    x_train.shape[0],
    IMAGE_WIDTH,
    IMAGE_HEIGHT,
    IMAGE_CHANNELS
)

x_test_with_channels = x_test.reshape(
    x_test.shape[0],
    IMAGE_WIDTH,
    IMAGE_HEIGHT,
    IMAGE_CHANNELS
)

print('x_train_with_channels:', x_train_with_channels.shape)
print('x_test_with_channels:', x_test_with_channels.shape)
```

↗ x\_train\_with\_channels: (60000, 28, 28, 1)  
x\_test\_with\_channels: (10000, 28, 28, 1)

## ✓ Normalize the data

Here we're just trying to move from values range of  $[0 \dots 255]$  to  $[0 \dots 1]$ .

```
x_train_normalized = x_train_with_chanel / 255
x_test_normalized = x_test_with_chanel / 255
```

```
# Let's check just one row from the 0th image to see color channel values after normalization.
x_train_normalized[0][18]
```

```
array([[0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.18039216],
       [0.50980392],
       [0.71764706],
       [0.99215686],
       [0.99215686],
       [0.81176471],
       [0.00784314],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ],
       [0.      ]])
```

## ✓ Build the model

We will use [Sequential](#) Keras model.

Then we will have two pairs of [Convolution2D](#) and [MaxPooling2D](#) layers. The MaxPooling layer acts as a sort of downsampling using max values in a region instead of averaging.

After that we will use [Flatten](#) layer to convert multidimensional parameters to vector.

The last layer will be a [Dense](#) layer with 10 [Softmax](#) outputs. The output represents the network guess. The 0-th output represents a probability that the input digit is 0, the 1-st output represents a probability that the input digit is 1 and so on...

```
model = tf.keras.models.Sequential()

model.add(tf.keras.layers.Convolution2D(
    input_shape=(IMAGE_WIDTH, IMAGE_HEIGHT, IMAGE_CHANNELS),
    kernel_size=5,
    filters=8,
    strides=1,
    activation=tf.keras.activations.relu,
    kernel_initializer=tf.keras.initializers.VarianceScaling()
))

model.add(tf.keras.layers.MaxPooling2D(
    pool_size=(2, 2),
    strides=(2, 2)
))

model.add(tf.keras.layers.Convolution2D(
    kernel_size=5,
    filters=16,
    strides=1,
    activation=tf.keras.activations.relu,
    kernel_initializer=tf.keras.initializers.VarianceScaling()
))

model.add(tf.keras.layers.MaxPooling2D(
    pool_size=(2, 2),
```

```

        strides=(2, 2)
    ))

    model.add(tf.keras.layers.Flatten())

    model.add(tf.keras.layers.Dense(
        units=128,
        activation=tf.keras.activations.relu
    ));

    model.add(tf.keras.layers.Dropout(0.2))


    model.add(tf.keras.layers.Dense(
        units=10,
        activation=tf.keras.activations.softmax,
        kernel_initializer=tf.keras.initializers.VarianceScaling()
    ))


```

 /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape` to `super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)`

Here is our model summary so far.

```
model.summary()
```

 **Model: "sequential"**

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 24, 24, 8)	208
max_pooling2d (MaxPooling2D)	(None, 12, 12, 8)	0
conv2d_1 (Conv2D)	(None, 8, 8, 16)	3,216
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 16)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 128)	32,896
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1,290

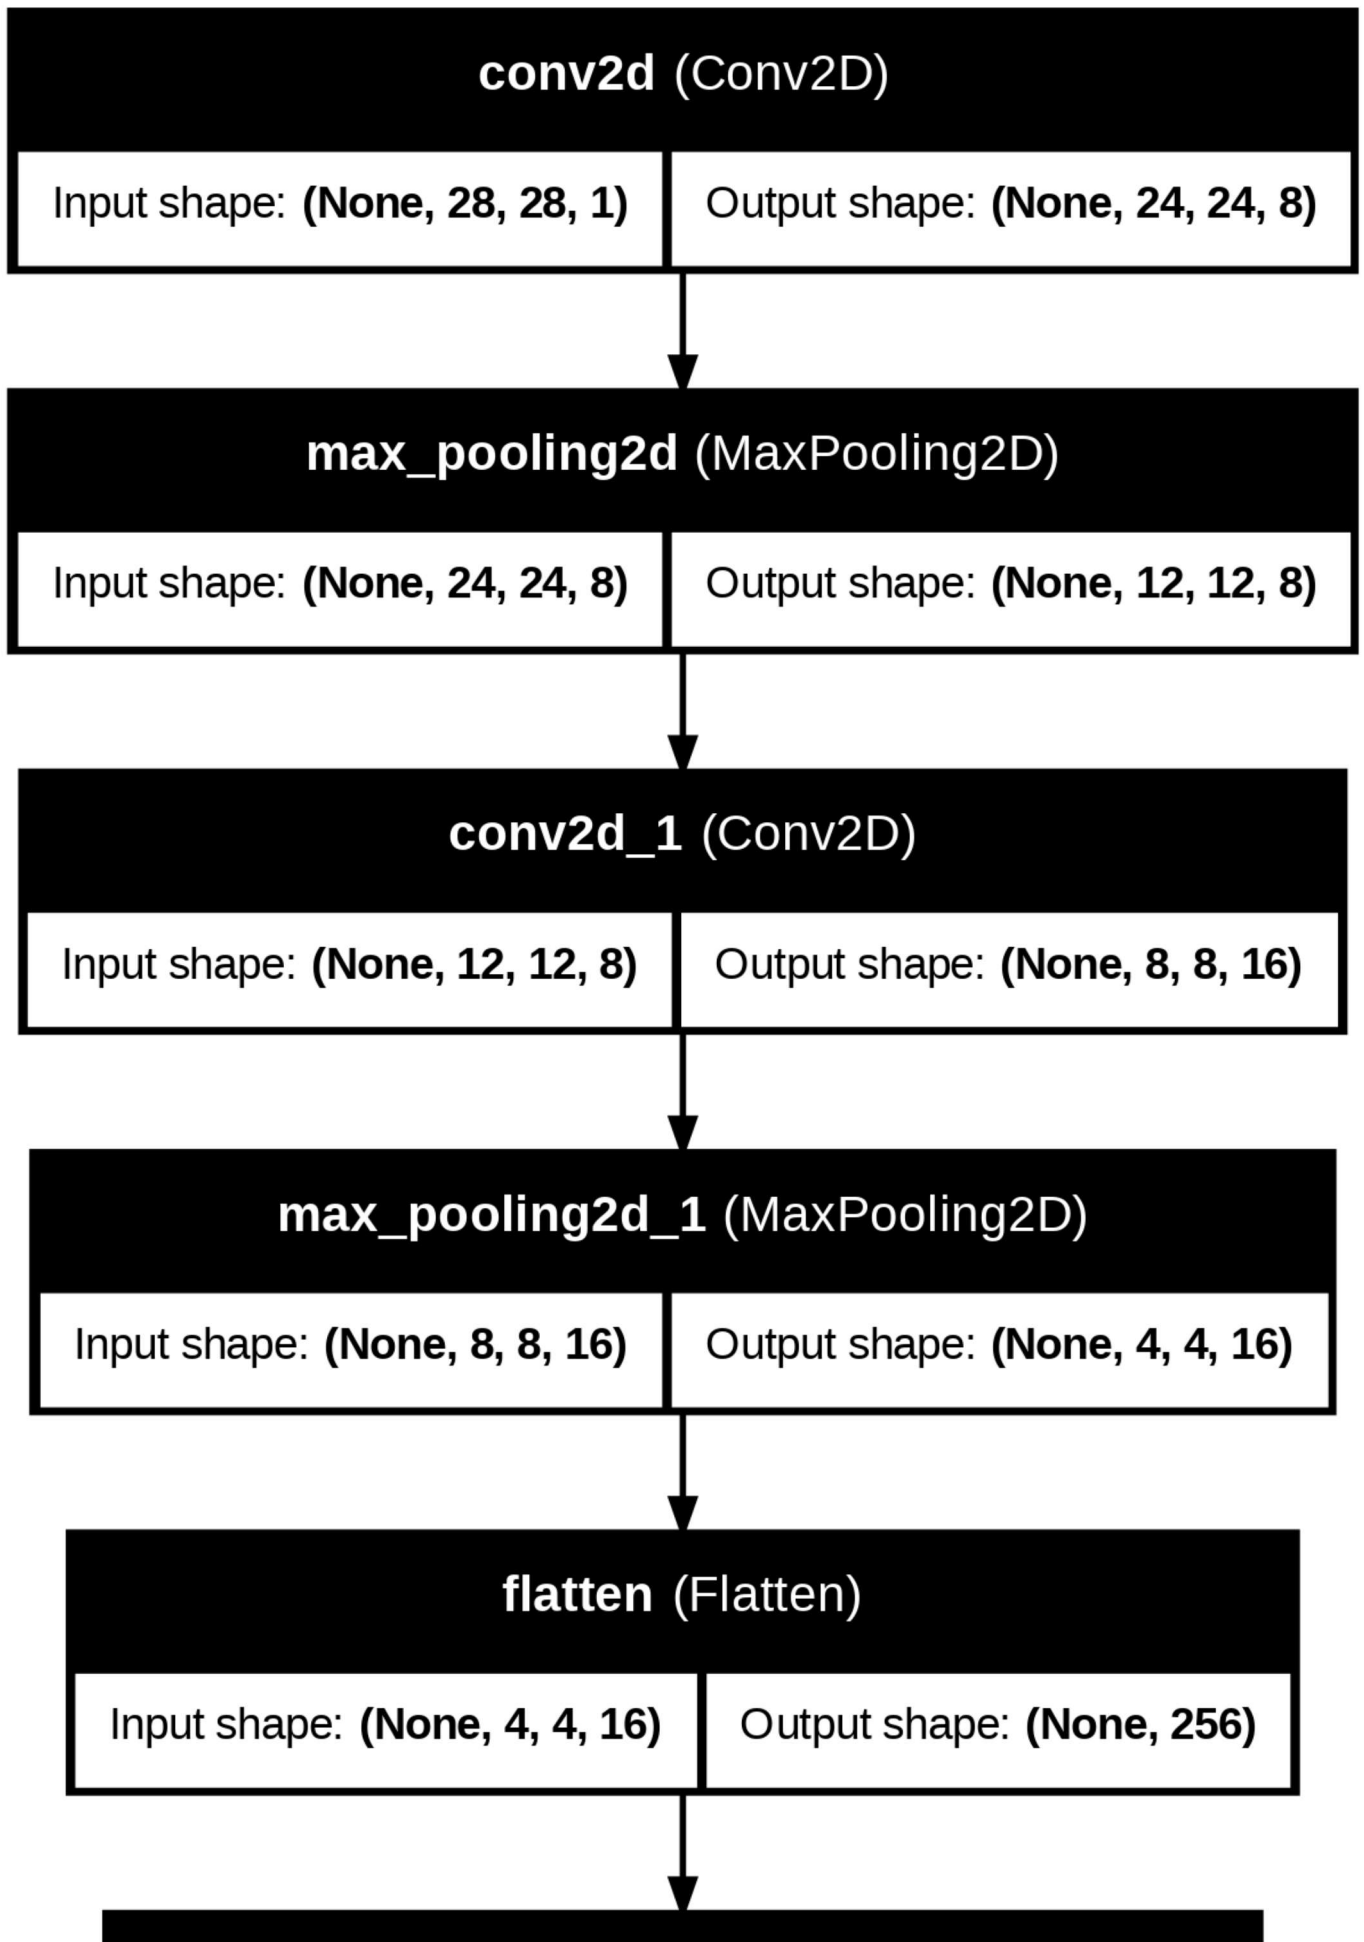
Total params: 37,610 (146.91 KB)  
 Trainable params: 37,610 (146.91 KB)  
 Non-trainable params: 0 (0.00 B)

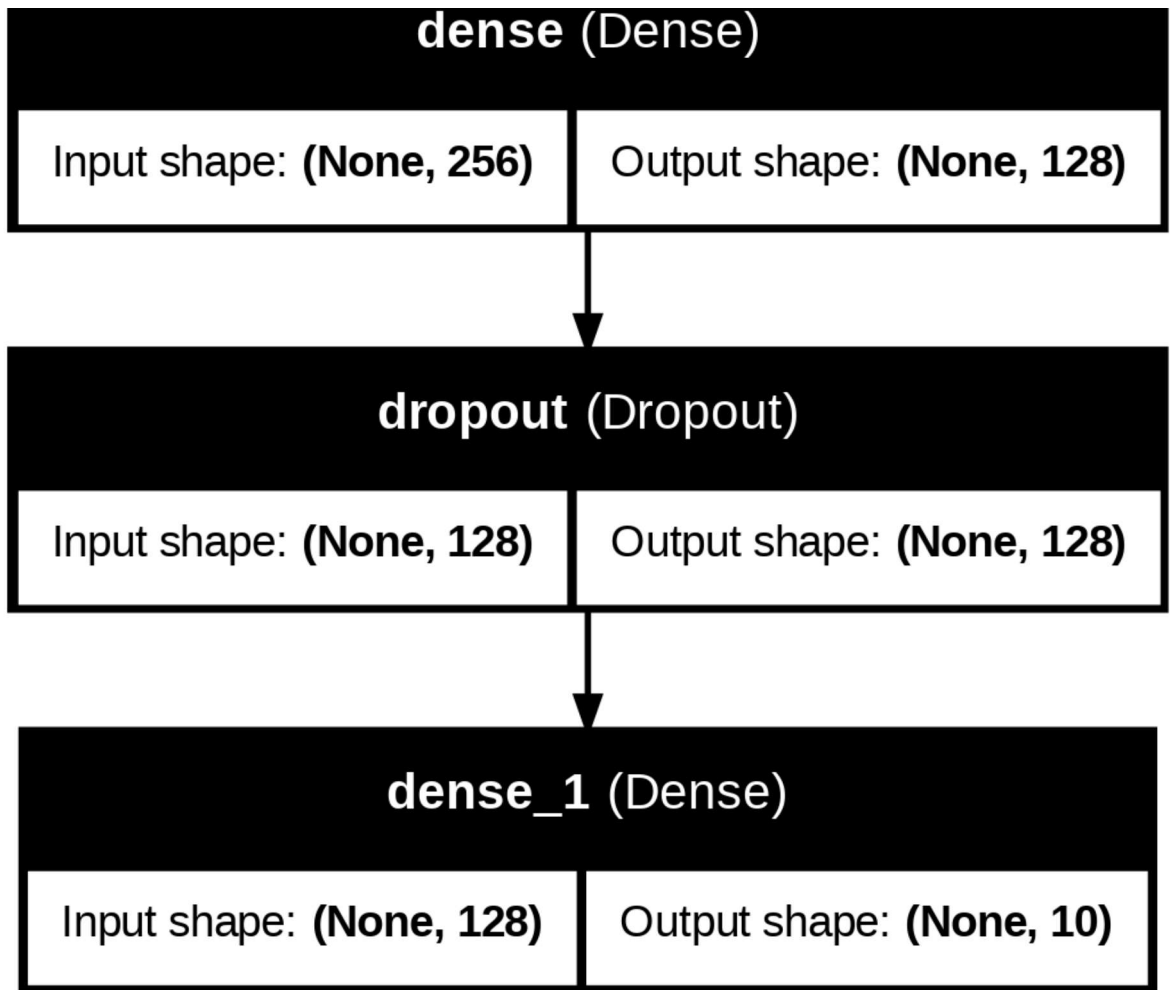
In order to plot the model the `graphviz` should be installed. For Mac OS it may be installed using `brew` like `brew install graphviz`.

```

tf.keras.utils.plot_model(
    model,
    show_shapes=True,
    show_layer_names=True,
)

```







## ✓ Compile the model

```
adam_optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)

model.compile(
    optimizer=adam_optimizer,
    loss=tf.keras.losses.sparse_categorical_crossentropy,
    metrics=['accuracy']
)
```

## ✓ Train the model

```
log_dir=".logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)
```


```
training_history = model.fit(
    x_train_normalized,
    y_train,
    epochs=10,
    validation_data=(x_test_normalized, y_test),
    callbacks=[tensorboard_callback]
)
```

```
↩ Epoch 1/10
1875/1875 ————— 39s 19ms/step - accuracy: 0.8495 - loss: 0.4769 - val_accuracy: 0.9834 - val_loss: 0.0538
Epoch 2/10
1875/1875 ————— 36s 17ms/step - accuracy: 0.9787 - loss: 0.0701 - val_accuracy: 0.9866 - val_loss: 0.0428
Epoch 3/10
1875/1875 ————— 39s 16ms/step - accuracy: 0.9855 - loss: 0.0481 - val_accuracy: 0.9875 - val_loss: 0.0376
Epoch 4/10
1875/1875 ————— 30s 16ms/step - accuracy: 0.9861 - loss: 0.0413 - val_accuracy: 0.9900 - val_loss: 0.0344
Epoch 5/10
1875/1875 ————— 41s 16ms/step - accuracy: 0.9897 - loss: 0.0319 - val_accuracy: 0.9902 - val_loss: 0.0315
Epoch 6/10
1875/1875 ————— 41s 16ms/step - accuracy: 0.9923 - loss: 0.0250 - val_accuracy: 0.9906 - val_loss: 0.0284
Epoch 7/10
1875/1875 ————— 30s 16ms/step - accuracy: 0.9919 - loss: 0.0235 - val_accuracy: 0.9899 - val_loss: 0.0335
Epoch 8/10
1875/1875 ————— 30s 16ms/step - accuracy: 0.9934 - loss: 0.0203 - val_accuracy: 0.9900 - val_loss: 0.0327
Epoch 9/10
1875/1875 ————— 43s 17ms/step - accuracy: 0.9935 - loss: 0.0206 - val_accuracy: 0.9914 - val_loss: 0.0335
Epoch 10/10
1875/1875 ————— 39s 16ms/step - accuracy: 0.9940 - loss: 0.0165 - val_accuracy: 0.9918 - val_loss: 0.0312
```

Let's see how the loss function was changing during the training. We expect it to get smaller and smaller on every next epoch.


```
plt.xlabel('Epoch Number')
plt.ylabel('Loss')
plt.plot(training_history.history['loss'], label='training set')
plt.plot(training_history.history['val_loss'], label='test set')
plt.legend()
```

 <matplotlib.legend.Legend at 0x7b0010769b90>



```
plt.xlabel('Epoch Number')
plt.ylabel('Accuracy')
plt.plot(training_history.history['accuracy'], label='training set')
plt.plot(training_history.history['val_accuracy'], label='test set')
plt.legend()
```

 <matplotlib.legend.Legend at 0x7b001071f5d0>



## ✓ Evaluate model accuracy

We need to compare the accuracy of our model on **training** set and on **test** set. We expect our model to perform similarly on both sets. If the performance on a test set will be poor comparing to a training set it would be an indicator for us that the model is overfitted and we have a "high variance" issue.

## ✓ Test set accuracy

```
%%capture
validation_loss, validation_accuracy = model.evaluate(x_test_normalized, y_test)

print('Validation loss: ', validation_loss)
print('Validation accuracy: ', validation_accuracy)
```



```
# Let's extract predictions with highest probabilities and detect what digits have been actually recognized.
predictions = np.argmax(predictions_one_hot, axis=1)
pd.DataFrame(predictions)
```

	0
0	7
1	2
2	1
3	0
4	4
...	...
9995	2
9996	3
9997	4
9998	5
9999	6

10000 rows × 1 columns


So our model is predicting that the first example from the test set is 7.

```
print(predictions[0])
```

```
7
```

Let's print the first image from a test set to see if model's prediction is correct.

```
plt.imshow(x_test_normalized[0].reshape((IMAGE_WIDTH, IMAGE_HEIGHT)), cmap=plt.cm.binary)
plt.show()
```



We see that our model made a correct prediction and it successfully recognized digit 7. Let's print some more test examples and correspondent predictions to see how model performs and where it does mistakes.

```
numbers_to_display = 196
num_cells = math.ceil(math.sqrt(numbers_to_display))
plt.figure(figsize=(15, 15))

for plot_index in range(numbers_to_display):
    predicted_label = predictions[plot_index]
    plt.xticks([])
    plt.yticks([])
```